

EVA ACQUISITION INTELLIGENCE · WHITEPAPER 01

The Manual Scanning Ends Here

A dataset-backed approach to acquisition deal sourcing —
built on a playbook and deal-outcome dataset no generic AI can match.

FOR PE FIRMS · SEARCHERS · ACQUISITION OPERATORS | JULY 2026

Executive summary

AI is already changing how deals are sourced and underwritten. McKinsey's 2026 Global Private Markets Report finds AI is improving the speed and consistency of sourcing, diligence, and portfolio monitoring, with leading GPs reporting productivity gains of 30 to 40 percent in analyst-intensive tasks.¹ EY-Parthenon's research shows GenAI can cut deal evaluation time by roughly 50 percent and enable 2.5× more deals to be evaluated through automated feasibility modelling.²

But McKinsey also flags an emerging and clear gap between the leading AI-forward sponsors and the rest — driven less by sector narratives and more by verifiable competitive moats, data advantages, embedded workflows, and execution capability.¹ The differentiator is no longer whether a firm uses AI. It is whether that AI is grounded in proprietary, outcome-validated data — or whether it is generic summarization that cannot survive an investment committee.


The gap Eva closes

Eva scans thousands of listings against your buy box and scores every deal against a proprietary deal-outcome dataset built from real, closed transactions — then hands you the three worth closing today. Every parameter is source-backed, so the output is defensible diligence for your IC and your LPs.

The problem: speed without defensibility

EY's 2026 Exit Readiness Study makes the distinction sharply: buyers are likely to distinguish between AI activity and AI strategy, and without clean, accessible, well-governed data, AI claims can quickly become difficult to evidence in diligence.³

Most AI deal-sourcing tools on the market today are AI activity, not AI strategy. They take a listing, summarize it, and produce a confident-looking memo. They are not trained on whether deals actually closed, what they closed at, or why they failed. The output is fast. It is not defensible. In a market where EY forecasts 2026 deal activity to surpass 2025 — driven by strategic AI-driven transformation⁴ — the cost of acting on undefendable speed is rising.



AI is already compressing deal cycles — but speed without ground truth is not yet a moat.

The shift: data advantages are the new moat

McKinsey is explicit that leading firms win on data advantages and embedded workflows rather than generic AI capability.¹ The most telling precedent is McKinsey's own Wave platform, which draws on a proprietary benchmark database of 500,000 individual initiatives to power value-creation underwriting.⁵

The principle

Proprietary, outcome-validated data is the only durable moat in AI-assisted dealmaking. Generic models can summarize any listing. They cannot tell you whether a deal like this one actually closed, at what multiple, and which parameters predicted the outcome — because that data does not exist in any public corpus.

How Eva works


Eva operationalizes a buy box as a scoring algorithm, then runs every new listing through it against the deal-outcome dataset.

- 1. Encode the buy box as parameters.** Your thesis becomes a weighted model — financial (price, SDE multiple, cash-flow margin), operational (owner-dependence, add-back quality, recurring revenue), structural (location, licensing, competitive density), and risk (concentration, lease, regulatory).
- 2. Score every listing against the outcome dataset.** Each listing is parsed, normalized, and scored — a comparison against the parameters of real deals that did and did not close, and at what terms.
- 3. Source-back every parameter.** No parameter is a guess. Every figure traces to the listing, a filing, a broker data field, or a public record. The memo carries its sources so it stands up to diligence.
- 4. Hand you the three worth closing today.** The scan produces a ranked list; the top three are packaged with the score, the source trail, and the specific next action. The long tail is logged for trend monitoring.



The moat: a proprietary deal-outcome dataset

The dataset is built from real acquisition activity — listings, LOIs, seller conversations, and closed/lost outcomes — each one feeding the scoring weights. Every deal Eva scores either validates or refutes a parameter, and the model recalibrates. This is the compounding loop: the more deals flow through, the sharper the scoring gets, and the wider the gap between Eva's output and any generic AI's summary.



A generic AI can replicate Eva's surface — the scanning, the memos, the formatting. It cannot replicate the dataset, because the dataset is the product of proprietary, ground-truth deal activity no public model can access.

What differentiates an Eva whitepaper

Most AI-in-M&A content is thought leadership — correct, generic, and unmoored from any specific deal. Eva whitepapers are different in one structural way: each one leaks a real, anonymized deal-outcome insight drawn from the dataset. The analysis is consultancy-grade, but it is anchored to a proprietary data point a generic AI cannot produce — a real scored deal, its parameters, and its outcome.

	Generic AI deal content	Eva whitepaper
Grounding	Public-web summarization	Proprietary deal-outcome dataset
Defensibility	Plausible, unsourced	Source-backed, IC-ready
Insight	General trend	Real anonymized scored deal + outcome
Moat	None (replicable)	Dataset compounds with every deal

Case study (anonymized, illustrative)

Senior-living acquisition · secondary market

A listing surfaced at a 4.1× SDE multiple. Manual scanning would have flagged it as borderline on cash-flow margin. Eva scored it 78/100: the margin cleared the dataset's threshold for closed deals in that asset class and geography, owner-dependence was low, and licensing density was favorable. The source trail confirmed add-backs and lease terms. The deal moved to LOI within the week. The scoring weights for that asset class were updated with the outcome — a dataset the next scan draws on.

Anonymized and illustrative of the scoring methodology; specific figures are representative.

The call to action

Eva does not replace your judgment. It replaces the weeks of manual scanning between a buy box and the three deals worth your attention — and it makes every score defensible to the people you answer to.

The manual scanning ends here.

Eva scans thousands of listings against your buy box — and hands you the three worth closing today, built on a playbook and deal-outcome dataset no generic AI can match.

Define your buy box and run your first scan: eva-acquisition.mangotec.ai

Sources

¹ McKinsey, Global Private Markets Report 2026 — <https://www.mckinsey.com/~media/mckinsey/industries/private%20equity%20and%20principal%20investors/our%20insights/mckinseys%20global%20private%20markets%20report/2026/global-private-markets-report-2026-full-report.pdf>

² EY-Parthenon-CREDAI, GenAI in Real Estate, June 2026 — https://www.ey.com/en_in/newsroom/2026/06/genai-could-deliver-up-to-50-percent-higher-sales-velocity-and-faster-launches-for-india-s-real-estate-sector-ey-parthenon-credai-report

³ EY, Global Private Equity Exit Readiness Study 2026 — https://www.ey.com/en_ly/insights/private-equity/private-equity-exit-readiness-study

⁴ EY, US M&A; Activity Insights 2026 — https://www.ey.com/en_us/newsroom

⁵ McKinsey, Unlocking Full Potential: Five Practices Reshaping PE Value Creation — <https://www.mckinsey.com/capabilities/transformation/our-insights/unlocking-full-potential-five-practices-reshaping-pe-value-creation>